

Intel Cloud Orchestration Networking

Winter Midterm Progress Report

Matthew Johnson, Cody Malick, and Garrett Smith

Team 51, Cloud Orchestra

Abstract

This document outlines the progress of the Cloud Orchestration Networking project over the Fall and Winter terms. It contains a short description of the project's purposes and goals, current progress, current issues, and any solutions to those issues. It also contains a week by week retrospective for all ten weeks of Fall term and the first half of Winter term.

CONTENTS

I	Project Goals	2
II	Purpose	2
III	Fall Progress	2
IV	Fall Term Week by Week Reports	3
IV-A	Background	3
IV-B	Weeks Zero Through Two	3
IV-C	Week Three	3
IV-D	Week Four	4
IV-E	Week Five	4
IV-F	Week Six	4
IV-G	Week Seven	4
IV-H	Week Eight	5
IV-I	Week Nine	5
IV-J	Week Ten	5
V	Fall Term Retrospective	6
VI	Winter Progress	6
VI-A	Network Testing Progress	7
VI-B	Development Progress	7
VII	Issues	7
VII-A	Environment Setup	7
VII-B	Network Testing	8
VIII	Winter Term Week by Week Reports	8
VIII-A	Winter Break and Week One	8
VIII-B	Week Two	8
VIII-C	Week Three	8
VIII-D	Week Four	9
VIII-E	Week Five	9
VIII-F	Week Six	9
IX	Winter Midterm Retrospective	10
	References	10

I. PROJECT GOALS

Our project is to first switch the Linux-created GRE tunnel implementation in Ciao to use GRE tunnels created by Open vSwitch. From that point we will switch the actual tunneling implementation from GRE to VxLAN/nvGRE based on performance measurements of each on data center networking cards. After this is completed, a stretch goal is to replace Linux bridges with Open vSwitch switch instances.

II. PURPOSE

The current implementation of Ciao tightly integrates software defined networking principles to leverage a limited local awareness of just enough of the global cloud's state. Tenant overlay networks are used to overcome traditional hardware networking challenges by using a distributed, stateless, self-configuring network topology running over dedicated network software appliances. This design is achieved using Linux-native Global Routing Encapsulation (GRE) tunnels and Linux bridges, and scales well in an environment of a few hundred nodes.

While this initial network implementation in Ciao satisfies current simple networking needs in Ciao, all innovation around software defined networks has shifted to the Open vSwitch (OVS) framework. Moving Ciao to OVS will allow leverage of packet acceleration frameworks like the Data Plane Development Kit (DPDK) as well as provide support for multiple tunneling protocols such as VxLAN and nvGRE. VxLAN and nvGRE are equal cost multipath routing (ECMP) friendly, which could increase network performance overall.

III. FALL PROGRESS

At present, the project is moving along smoothly. Our testing environment has been set up and is networked appropriately. Each Intel NUC (Next Unit of Computing) has Clear Linux installed. Come Winter term, we will get Ciao set up on each machine and begin development on Ciao. Software development on the project has yet to begin as we have just wrapped up the design phase.

Designing has been quite helpful in developing our understanding of the project, its goals, and purpose. Because this is a small component of a very complicated system, taking the time to investigate what a software defined network is, why it's being used, and why we are implementing the piece that we are has been quite beneficial.

While in the design phase, we found an extremely useful library for interfacing with Open vSwitch, libovsdb. Libovsdb is a library written in the Go programming language that allows for simple and efficient calls to the OVS Database Management Protocol. Interfacing with OVS is going to be a very large portion of the project for us, so finding the library is quite the boon. Here is a quick example of how this library functions:[1]

Listing 1. Example insert operation using libovsdb

```
// simple insert operation
insertOp := libovsdb.Operation{
```

```

Op: "insert",
Table: "Bridge",
Row: bridge,
UUIDName: namedUUID,
}

```

The above example can be reused for all major operations in the OVS Database Management Protocol. Other example operations include select, delete, and update. Using just the operations listed here, we can accomplish most of the needed configuration changes within Open vSwitch.

IV. FALL TERM WEEK BY WEEK REPORTS

A. Background

Over the Summer, 2016, Matthew worked as a Software Engineering Intern for the Advanced Systems Engineering (ASE) group within the Open Source Technology Center (OTC) that is in turn within the Software Services Group (SSG) at Intel. Matthew's coworkers had a need to integrate Open vSwitch in their cloud orchestration software (Cloud Integrated Advanced Orchestrator, or Ciao) but did not have the man hours to contribute time to it. Matthew worked with the team to propose the project as a Senior Capstone project at Oregon State University.

Matthew identified two other students, Cody Malick and Garrett Smith, as intelligent hard workers who would benefit the project. Because of this Robert Nesius, the Intel Engineering Manager serving as our client, requested Matthew, Cody, and Garrett specifically for this project.

B. Weeks Zero Through Two

During the first week of class, we all visited Intel in Hillsboro. The principal engineer in charge of Ciao networking, Manohar Castelino, gave us all an explanation of Ciao, how the networking works, and what he expects us to accomplish.

It was also during this time that Rob provided us with five Intel NUCs that would serve as our local cluster. We found out that we needed to register the MAC addresses with the university, and communicated this need to Todd Shechter, the Oregon State University Director of Information Technology.

C. Week Three

During week three we attempted to install Clear Linux OS for Intel Architecture [2] on all five Intel NUCs. We were unsuccessful because the installer requires a network connection to download the packaging. At this point, network access had yet to be approved by OSU IT.

Network access is required to install Clear Linux, Ciao, and access the machines remotely. Since Clear Linux is a datacenter OS, not a desktop OS, it does not support wireless internet connections. Because ethernet is required our hardware must be

registered with the university. If we were unable to obtain network access for the hardware on OSU's network we would have needed to find somewhere else to house it.

We also wrote our problem statement in week three, earlier than most groups were able to, since we were ahead of schedule with regard to meeting our team and choosing a project.

D. Week Four

During week four Matthew installed Clear Linux on the Intel NUCs. Since there the networking issues had still not been solved he brought them to his house to use the wired connection there. At this point our hardware had been registered with the university, but for unknown reasons our NUCs were not connecting to the network. Todd Shechter was devoting a lot of his time to help us debug, but we were not yet successful. He set us up with two five-port switches in Kevin McGrath's lab, but they were not receiving IP addresses from the network.

We had our problem statement reviewed and were waiting for feedback by the end of the week. We also started working on the client requirements document, though much of our contributions were simple outline and templating work.

E. Week Five

This week we spent most of our time writing the rough draft for the requirements document. We turned it in by the end of the week and were satisfied with our progress. The final draft of the requirements document was due the next week, week six. This week we also turned in our signed copy of the problem statement.

This week our networking issues were resolved. We had another email conversation with Todd Shechter, who was at a loss as to why we could not access the network from the NUCs. He granted us access to an HP switch we had successfully connected via in the past. On Thursday, we went to move all our NUCs to the new switch, but tried out the network on the mini switches one last time. This time they all worked. Our hardware was now set up and ready to go.

F. Week Six

This week we finished writing the requirements document that was due at the end of the week. After the rough draft we turned in the previous week we continued working on it ourselves until Tuesday. On Wednesday Frank emailed us some suggestions and we addressed those right away. We got the document signed that afternoon by Rob Nesius, our client at Intel.

G. Week Seven

This week we started working on the technical review document due the following Monday. This document outlines nine different components of our system. For each component we explored three different technologies that could be used to implement the component. Since our project is implementing a component of a larger system, it was difficult for us to come

up with nine components and three technologies each (twenty-seven different options). We spent much of our week working together to figure out how to split the project up.

H. Week Eight

This week we submitted the tech review document. It was a lot of work, but we got it in on time unlike many other capstone groups. Garrett researched software switch options, network latency tools, and network throughput tools. Matthew dealt with high-level language, testing, and logging tools. Cody handled the network-specific implementation pieces, such as packet protocols, network virtualization implementation, and bridge implementation.

I. Week Nine

This was a short week with Thursday and Friday given over to the Thanksgiving holiday. We started working on the design document but did not make much headway before breaking for the holiday weekend.

Our team also got together and talked about how we were going to execute the final presentation video for a few minutes this week.

J. Week Ten

This week we focused on the design document due Friday. We spent time researching design strategies and writing up our plan to execute. During this research we found a very useful Go library that interfaces with Open vSwitch. This library will simplify our implementation, allowing us more time to do network performance testing, which the client is very interested in. Our client signed the document with a half-hour to spare before the deadline, and the technical advisors for the project, Manohar Castelino and Tim Pepper, indicated they were impressed with the design we had outlined.

V. FALL TERM RETROSPECTIVE

Week	Positives	Deltas	Actions
0-2	Met the Intel team, studied project goals, purposes	Write project abstract	Research project, write project abstract
3	Started testing hardware setup, problem statement first draft submitted	Resolve networking issues	Contact Todd to get NUC network authorization, write final draft of problem statement
4	Completed hardware setup	Get problem statement signed	Email project owners and get problem statement approved
5	Problem statement submitted, completed first draft of requirements document	Submit final draft of requirements document	Update requirements document, email project owners, get approval via signature
6	Final draft of requirements document submitted	Begin work on tech review	Research technologies for tech review
7	First draft of tech review completed	Finalize tech review	Update and submit tech review
8	Submitted tech review	Begin design document	Research project design steps and implementation details
9	Began work on design document	Complete design document	Fill out the rest of the design document over Thanksgiving weekend
10	Completed design document, began work on final report and final presentation	Complete final report, complete final presentation	Over the weekend, complete final report, create slides for final presentation

VI. WINTER PROGRESS

Winter term has started off a little bumpy for the team. We spent the first four weeks focusing on setting up our development environment, as well as ironing out issues related to network testing. These topics will be covered in-depth in the following 'issues' section. Our current progress, as far as development is concerned, is working on the first feature, implementing OVS generated GRE tunnels. While we originally planned to have that feature completed by the end of the fifth week, we are in contact with our friends at Intel to figure out some of the issues we've been running into. Good progress on testing the basic NUC setup, and will be covered in the following section. Lastly, our current development progress, and issues related to deployment of OVS in Ciao will be discussed.

A. Network Testing Progress

B. Development Progress

VII. ISSUES

In this section, we will cover the issues we ran into at the beginning of term, as well as the solutions, and discussions that were had around them. The first two major issues are related to environment setup. These two issues consumed a large portion of the first part of the term.

A. Environment Setup

The first, and arguably most major issue that was experienced, was setting up the Ciao cluster. The first step of the project is, quite simply, getting Ciao working in the state that it was provided to us. This ultimately should have been fairly straightforward, but we ran into two major issues. The first of which was the management of certificates in Clear Linux.

Clear Linux, Intel's custom Linux distribution, is a fast OS designed to take advantage of Intel CPU's advanced features that often go unutilized. This was an obvious choice for our team. We initially set up all the NUCs with Clear Linux installed, and proceeded with deployment. We spent about a week tracking down an issue with Clear Linux. Specifically, that Clear Linux manually manages the network trust store. This isn't an issue, but the behavior of Clear Linux was slightly different than Ubuntu's. After some investigation, we reported the bug to the Clear Linux team, who had a bug fix shipped in the next release.

The second major issue we encountered was that of getting FQDN's in Ciao, or Fully Qualified Domain Names. These were important to the controller node of Ciao as it needed to identify the compute nodes needed to deploy software onto. As we were trying to deploy, we found that Ciao kept attempting to connect to the hostname of the device instead of the FQDN. For example, the control node would try to connect to `fw-dear205-ciao-nuc0` instead of the FQDN, `fw-dear205-ciao-nuc0.engr.oregonstate.edu`. A screenshot of what is expected of the python call is show below. At first, we thought this issue was because of caching in the OS, but quickly found out this was not the case after manually clearing the cache, as well as rebooting the system. After some further investigation, we found that Python3 was only getting the hostname on Clear Linux instead of the FQDN. We reported this problem to the Clear Linux team as well, and they have since shipped a fix.

Fig. 1. Expected result of Python FQDN call

```
garrett@fw-dear205-ciao-nuc4:~$ python3
Python 3.5.2 (default, Nov 17 2016, 17:05:23)
[GCC 5.4.0 20160609] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import socket
>>> socket.getfqdn()
'fw-dear205-ciao-nuc4.engr.oregonstate.edu'
>>>
```


We found out later that the Ciao dev team usually uses Ubuntu as its development environment. With that in mind, we decided to move to an Ubuntu development environment to make it as consistent as possible. We have since then installed Ubuntu 16.10 on all the machines. Also at the recommendation of the Ciao team, we started using Ciao-Down, a single virtual machine development environment. With this set up, we were able to quickly get to work on the development of the first feature. While we are not currently using a fully deployed version of Ciao for development, we will work on getting Ciao tested with the new features in our full five NUC setup once we've made progress on development.

B. Network Testing

VIII. WINTER TERM WEEK BY WEEK REPORTS

A. Winter Break and Week One

Over Christmas break Cody and Matthew tried to get Ciao set up on the cluster using a manual installation method. We ran into some tough issues and communicated them with the Intel team. The Intel team pointed us to a much more recently updated setup document that included automated deployment.

During week one there was still an issue in the version of Clear Linux (our target distribution) that we were running that caused docker certification to be broken. This was fixed within the week and we continued trying to set up Ciao.

B. Week Two

Garrett wrote several scripts to collect network performance data between nodes. He made significant progress in parsing bandwidth data into csv format for graphical representation.

This week we continued to work on the Ciao deployment via automated ansible playbooks. We encountered several issues from the start and worked through them one by one. Issues included errors in the ansible playbooks regarding yaml parsing and fqdn configuration. This type of issue was resolved by hardcoding the playbooks for our specific setup. By the end of the week we were seeing certification management issues in the build.

C. Week Three

This week we believed we were successful in deploying Ciao and began implementation of Open vSwitch components. Garrett has begun working out network measurements for our initial benchmarks. Garrett discovered that the network was not being set up properly, however, due to issues with the OSU network.

The first half of the week was spent debugging our Ciao deploy with members of the Ciao development team at Intel. After several email conversations and debug steps, we were advised to switch our operating system to Ubuntu because of various certification issues in Clear Linux. This, along with running from within a ciao-deploy docker container, fixed our issues and allowed us to properly initialize the cluster. As mentioned earlier, our network was not properly set up.

D. Week Four

During week four we continued tentative development on the OVS modules for Ciao. Some of the necessary functions have been written, but it has been difficult to test this over the physical cluster. Garrett, who has worked on the OSU network in the past, thinks the network may not be assigning IPs to the cluster. The OSU network DHCP servers will not assign an IP unless the MAC address has been registered with the university. To get around this we may have to do single-vm setup or find a way to set up the cluster independent from the OSU network.

Ciao provides `ciao-down`, a tool that helps set up a single-vm environment for testing. This will be helpful once we get it spun up.

We were still having issues gathering initial network metrics due to inaccessible nodes on our physical cluster. Garrett is working through this but the best way to address this may be single-vm for now.

E. Week Five

Early in the week we got `ciao-down` (the single-vm setup for Ciao) working to test our code.

We have a schedule we are following, and our initial goal was to complete the OVS module by the end of week five. Although we had a module written of something we were hoping might work, we were unable to get it to integrate properly with the rest of Ciao. It is likely we were misunderstanding some things in the calling hierarchy. Matthew communicated the issues we were seeing with Manohar Castelino, one of our clients and a software defined networking expert. This led to a conversation about whether or not it was strictly necessary to create the bridges in Open vSwitch or if we could simply create the GRE tunnels themselves.

As far as the physical cluster goes, we realized that the OSU network will not lease IP addresses unless the MAC address is registered with the University. Since Ciao uses the network's DHCP server to assign IP addresses to the nodes this is a big issue. Garrett has been setting up the physical cluster with a DHCP server running on the switch, with only our deployment NUC connected to the internet. He has been running into issues setting up the Keystone server once he took our controller NUC off the internet. He was trying to modify ansible tasks to get around the issue when I left campus today. If he is not able to get this to work we will probably ask the university for a subnet we can run a non-MAC-locked DHCP server on, but Garrett has worked in OSU IT before and said we are unlikely to get permission for that. The number of errors the OSU network has generated for us has been a little frustrating.

F. Week Six

Matthew met with Manohar in-person at Intel on Tuesday to discuss the issues with OVS bridges. It turns out Open vSwitch will not attach tunnels to non-OVS bridges. This expands our scope somewhat to build a full OVS framework for Ciao, instead of just OVS tunnels.

The rest of the week was devoted to working on the midterm progress report for Winter term.

IX. WINTER MIDTERM RETROSPECTIVE

Week	Positives	Deltas	Actions
1	Learned about automated deployment for Ciao	Worked through several deployment issues	Attempted to deploy Ciao
2	Scripts to collect network statistics completed	Resolved old issues in deployment but uncovered more	Hardcoding variables in the ansible deployment resolved some new issues
3	Successfully initiated the cluster	Switched OS on the cluster to Ubuntu	Started working in single-vm mode
4	Wrote some of the module for OVS tunnels	OSU network DHCP will not assign IPs to unregistered MACs	Developed in ciao-down
5	OVS module mostly written, but not integrated	Discussed whether OVS bridges were required for OVS tunnels	Tried to set up local DHCP server on switch
6	Met with Manohar regarding OVS bridges	OVS bridges are required for OVS tunnels	Worked on progress report for midterm

REFERENCES

- [1] Socketplane. (2016, dec) play_with_ovs.go. [Online]. Available: https://github.com/socketplane/libovsdb/blob/example/play_with_ovs.go
- [2] Intel. (2016, dec) clearlinux.org. [Online]. Available: <https://clearlinux.org/>